

# CSIT 231: Systems Programming

Michelle Zhu

Computer Science Department  
Montclair State University



- Slides revised from CS 252 Systems Programming
  - Gustavo Rodriguez-Rivera
    - Computer Science Department
      - Purdue University



---

# The UNIX Operating System

---

# What is an Operating System

---

- # An Operating System (OS) is a program that sits in between the hardware and the user programs.
  - # It provides:
    - Multitasking - Multiple processes running in the same computer
    - Multiuser - Multiple users using the same computer
    - File system – Storage
    - Networking – Access to the network and internet
-

# What is an Operating System

---

- Window System – Graphical use interface
  - Standard Programs – Programs such as a web browser, task manager, editors, compilers etc.
  - Common Libraries – Libraries common to all programs running in the computer such as math library, string library, window library, c library etc.
  - It has to do all of the above in a secure and reliable manner.
-

# A Tour of UNIX

---

- # We will start by describing the UNIX operating system (OS).
  - # Understanding one instance of an Operating System will help us understand other OSs such as Windows, Mac OS, Linux etc.
  - # UNIX is an operating system created in 1969 by Ken Thompson, Dennis Ritchie, Brian Kernighan, and others at AT&T Bell Labs.
  - # UNIX was a successor of another OS called MULTICS that was more innovative but it had many problems.
  - # UNIX was smaller, faster, and more reliable than MULTICS.
-

# A Tour of UNIX

---

- # UNIX was initially created to support typesetting (edition of documents).
  - # By having the programmers being the users themselves of the OS. UNIX became the robust, practical system that we know today.
  - # UNIX was written in “C” (95%) and assembly language (5%).
  - # This allowed UNIX to be ported to other machines besides Digital Equipment (DEC)’s PDP11.
-

# BSD UNIX

---

- # UNIX was a success in the universities.
- # Universities wanted to modify the UNIX sources for experimentation. Berkeley created its own version of UNIX called BSD- UNIX.
- # POSIX (The Portable Operating System Interface) is a family of standards specified by the IEEE Computer Society for maintaining compatibility between operating systems
- # Sockets, FTP ( File Transfer Protocol) which is is a standard network protocol used for the transfer of computer files from a server to a client on a computer network., and Mail etc came from BSD UNIX.





---

# The UNIX File System

---

# UNIX File System

---

- # UNIX has a hierarchical File System

- # Important directories

  - / Root Directory

  - /etc OS Configuration files

    - /etc/passwd – User information

    - /etc/groups – Group information

    - /etc/inetd.conf – Configuration of Internet

      - inetd**: internet service daemon

    - /etc/rc.\*/ - OS initialization scripts for  
different services

- # Deamons – Programs running in the background implementing a service. For each configured service, it listens for requests from connecting clients.

---

# UNIX File System

---

`/dev` – List of devices attached to the computer

`/usr` – Libraries and tools

`/usr/bin` – Application programs such as `grep`, `ls` et

`/usr/lib` – Libraries used by the application programs

`/usr/include` – Include files (`.h`) for the libraries

`/home` – Home directories

---

# Users

```
$ cat /etc/passwd
```

Usually, the first line describes the root user, followed by the system and normal user accounts. New entries are appended at the end of the file.

Each line of the `/etc/passwd` file contains seven comma-separated fields:

Output

```
mark:x:1001:1001:mark,,,:/home/mark:/bin/bash
[---] - [---] [---] [-----] [-----] [-----]
|      |      |      |      |      |      |
|      |      |      |      |      |      +--> 7. Login shell
|      |      |      |      |      +-----> 6. Home directory
|      |      |      +-----> 5. GECOS
|      |      +-----> 4. GID
|      +-----> 3. UID
|      +-----> 2. Password
+-----> 1. Username
```

UNIX was designed as a multiuser system.

The database of users is in `/etc/passwd`

The encrypted password used to be stored also here. Now it is stored in `/etc/shadow`

# Users

---

## # Commands for users

- adduser – Adds a new user
- passwd – Change password.

## # There exist a special user called “root” with special privileges.

## # Only root can modify files anywhere in the system.

## # To login as root (superuser) use the command

## # “su”.

Only root can add users or reset passwords.

---

# Groups

---

- # A “group” represents a group of users.
  - # A user can belong to several groups.
  - # The file `/etc/group` describes the different groups in the system.
-

# File Systems

---

# The storage can be classified from fastest to slowest in the following

- Registers
  - Cache
  - RAM
  - Flash Memory
  - Disk
  - CD/DVD
  - Tape
  - Network storage
-

# Disk File Systems


---

- # A file system is a logical collection of files on a partition or disk.
    - Unix divided physical disks into logical disks called partitions
    - A partition is a container for information
  - # Everything in Unix is considered to be a file, including documents, hard-drives, modems, keyboards, printers and even some inter-process and network communications as simple streams of bytes exposed through the filesystem
-




# Disk File Systems

The disk is a an electromagnetic and mechanical device that is used to store information permanently.



# Disk File Systems



A Block is the intersection between a sector and a track

Disks when formatted are divided into **sectors, tracks and blocks**.

Disks are logically divided into partitions.

A partition is a group of blocks.

Each partition is a different file system.

---

# Disk File Systems

- # The command **df -h** (disk free) displays the **disk space usage in human readable**.

```
$ df --human-readable
Filesystem      Size  Used Avail Use% Mounted on
/dev/root        110G   45G   61G   43% /
devtmpfs         12G     0    12G    0% /dev
tmpfs            12G   848K   12G    1% /run
/dev/sda1        1.6T   1.3T   191G   87% /home
/dev/sdb1        917G   184G   687G   22% /penguin
/dev/sdc1        57G    50G   4.5G   92% /sneaker
/dev/sdd1        3.7T   2.4T   1.3T   65% /tux
```

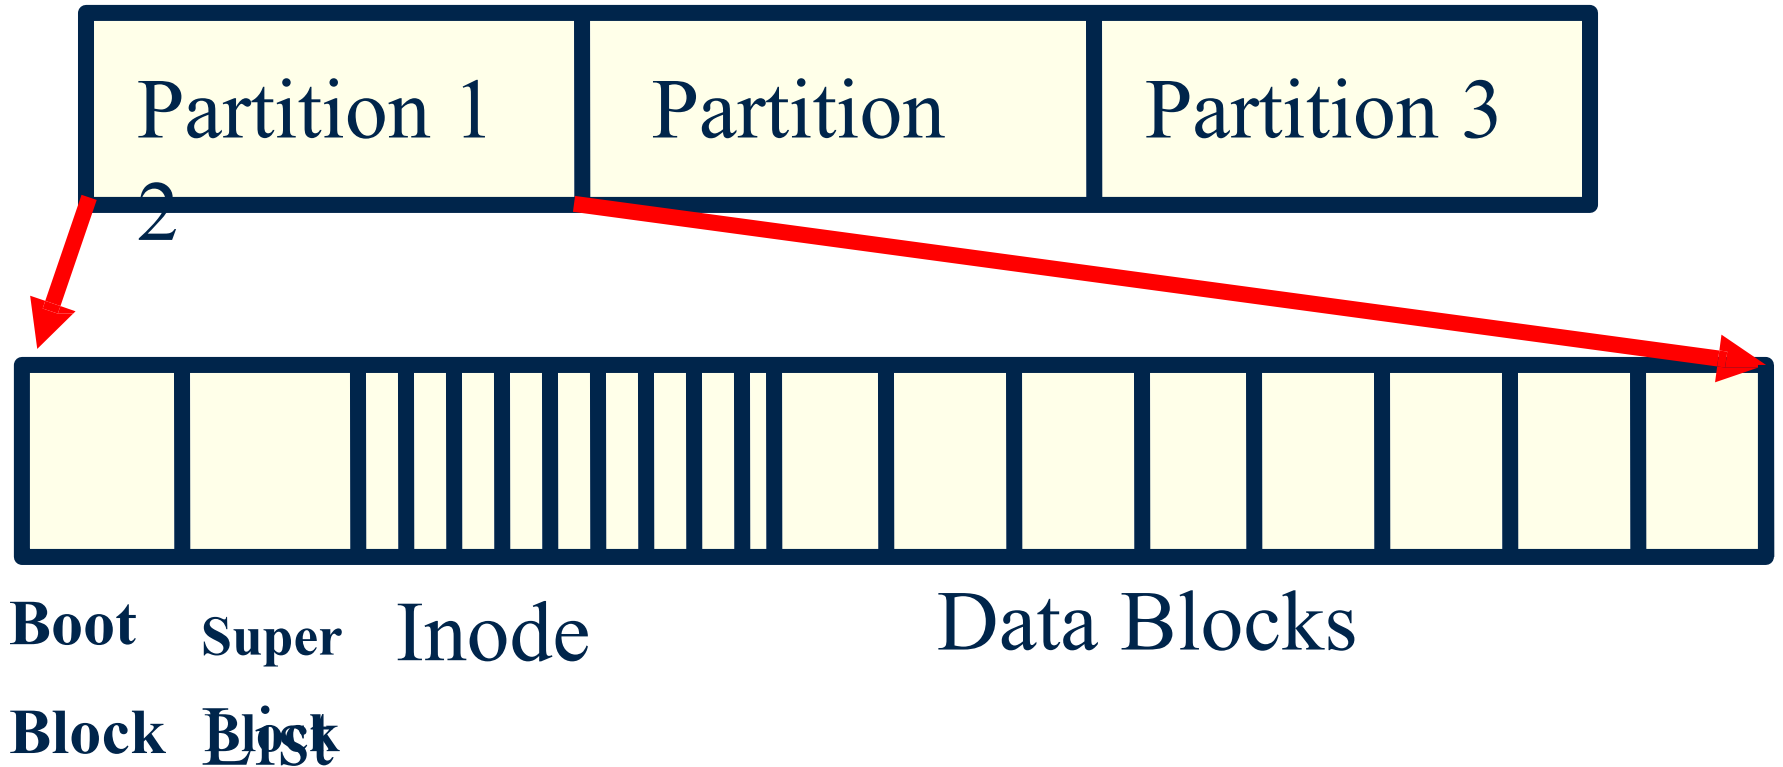
- # The **du** (disk usage) command enables you to specify directories to show disk space usage on a particular directory in blocks.

```
$du /etc
10      /etc/cron.d
126     /etc/default
6       /etc/dfs
...
$
```

The **-h** option makes the output easier to comprehend –

```
$du -h /etc
5k      /etc/cron.d
63k     /etc/default
3k      /etc/dfs
...
$
```

# Disk File System



# Disk File System

---

- # Each partition is divided into:
  - **Boot Block** – Located in the first few sectors of a file system. contains the initial bootstrap program used to load the operating system.
  - **Superblock** – Describes the state of the file system: the total size of the partition, the block size, pointers to a list of free blocks, the inode number of the root directory.
  - **Inode-list** – It is a linear array of I-nodes. One to one mapping of a file to an Inode. An inode has information about a file and what blocks make the file.
  - **Data Blocks** – Store the file data.

Thus, while users think of files in terms of file names, Unix thinks of files in terms of inodes.

---

# File and Inode

- *You store your information in a file, and the operating system stores the information about a file in an inode (sometimes called as an inode number).*
- Whenever a user or a program needs access to a file, the operating system first searches for the exact and unique inode (inode number), in a table called as an inode list/table. In fact the program or the user who needs access to a file, reaches the file with the help of the inode number found from the inode list/table.

# I-node information

- An i-node represents **a file** in disk. Each i-node contains **meta info** of the file:
  1. Flag/Mode
    1. Read, Write, Execute (for Owner/Group/All) RWX RWX  
RWX Also tells if file is directory, device, symbolic link
  2. Owners
    1. Userid, Groupid
  3. Time Stamps
    1. Creation time, Access Time, Modification Time.
  4. Size
    1. Size of file in bytes
  5. Ref. Count –
    1. Reference count with the number of times the i-node appears in a directory (hard links).
    2. Increases every time file is added to a directory. The file the i-node represents will be removed when the reference count reaches 0.


# I-node information

---

- # The I-node also contains block index with the blocks that form the file.
  - # To save space, the block index uses indices of different levels.
  - # This benefits small files since they form the largest percentage of files.
  - # Small files only uses the direct and single-indirect blocks. When files get bigger, it will use more and more indirect
  - # block.  
This saves in space spent in block indices.
-



# Inode as a Data Structure



## **Mode:**

This keeps information about two things, one is the permission information, the other is the type of inode, for example an inode can be of a file, directory or a block device etc.

**Owner Info:** Access details like owner of the file, group of the file etc.

**Size:** This location store the size of the file in terms of bytes.

**Time Stamps:** it stores the inode creation time, modification time, etc.

Now comes the important thing to understand about how a file is saved in a partition with the help of an inode.

**Block Size:** Whenever a partition is formatted with a file system. It normally gets formatted with a default block size. Now block size is the size of chunks in which data will be spread. So if the block size is 4K, then for a file of 15K it will take 4 blocks(because  $4K \times 4 = 16K$ ), and technically speaking you waste 1 K.

<https://www.slashroot.in/inode-and-its-structure-linux>

Everything except **file name and data** are stored here

# I-node information

---

- Modern File System has 15 pointers for each I-node structure

Direct block –

- Points directly to the block. There are **12** of them in the structure

- Single indirect –

- Points to a block table that has various entry's, say 128 or even more, There are **1** of them.

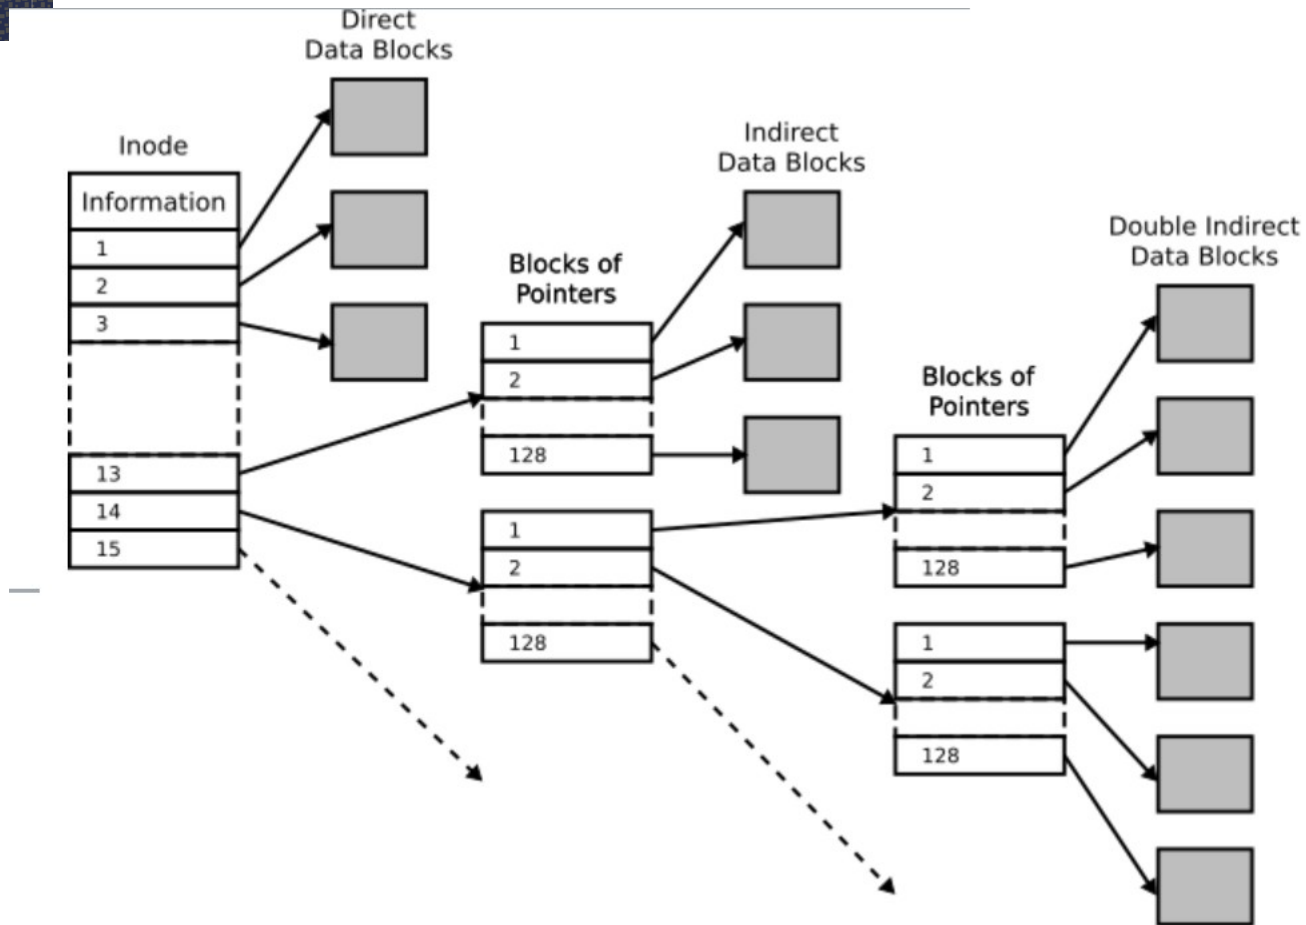
- Double indirect –

- Points to a block table of 128 entries which then points to another block table of 128. There is **1** of them.

- Triple Indirect -

- Points to a block table of 128 entries which then points to another block table of 128 that points to another table of 128. There is **1** of them.
-

# I-node Structure



# I-node information

---

■ Assume 1KB for each block and 128 entries in each block table.

■ Direct block =  $12 * 1Kb = 12Kb$

■ Single indirect =  $1 * 128 * 1Kb = 128 Kb$

■ Double indirect =  $1 * 128 * 128 * 1Kb = 16 Mb$

■ Triple indirect =  $1 * 128 * 128 * 128 * 1Kb$   
= 2 Gb

More and more block pointers will be used when the file gets bigger.

---

# I-node information

---

- # Most of the files in a system are small.
  - # This also saves disk access time since small files need only **direct blocks**.
    - 1 disk access for the I-Node
    - 1 disk access for the datablock.
  - # An alternative to the multi-level block index is a ***linked list***. Every block will contain a pointer to the next block and so on.
  - # Linked lists are slow for random access.
-

# Directory Representation and Hard Links

- # Directories are just tables that associate file names with inode numbers. Each file name and inode number pair in a directory is called a link. Multiple names for the same inode are called “hard links”.
- # Each file name is mapped to only one single inode number, but one file inode number may have many names that map to it.
- # An I-node may appear in multiple directories.
- # The reference count in the I-node keeps track of the number of directories where the I-node appears.
- # When the reference-count reaches 0 by run rm command, the inode has no name and is freed.

```
$ ls -li /usr/bin/perl*
266327 /usr/bin/perl          266329 /usr/bin/perl.doc.stub
266327 /usr/bin/perl5.14.2    266330 /usr/bin/perl.vp
266331 /usr/bin/perlbug       266331 /usr/bin/perlthanks
266328 /usr/bin/perl.doc
```

# Hard Links

---

- # In some OSs, the reference count is incremented when the file is open.
- # This prevents the file from being removed while it is in use.
- # Hard Links has the same inode value as the original. If the original inode is updated, all hard links are updated.  
It cannot cross partitions, that is, a directory cannot list an I-node of a different partition.

Example. Creating a hard link to a target-file in the current directory

***ln target-file name-link***

---



Example hard link versus soft/symbolic link

<https://linuxgazette.net/105/pitcher.html>

# Soft-Links

- ✦ Directories may also contain Soft-Links, which acts like shortcut in Windows.

- ✦ A soft-link is a pair of the form  
(file name, i-node number-with-file-storing-path)

Where path may be an absolute or relative path in this or another partition.

- ✦

Soft-links can point to files in different partitions. A soft-link/symbolic link is referring to the original file and not its inode value, then replacing into another folder/removing the original file will break the symbolic link, or create a dangling link.

Example:

***ln -s target-file name-link***



# File Ownership

---

- # The Group Id and owner's User ID are stored as part of the file information
  - # Also the creation, modification, and access time are stored in the file in addition to the file size.
  - # The time stamps are stored in seconds after the Epoch (0:00, January 1<sup>st</sup>, 1970).
-

# File Permissions

# The permissions of a file in UNIX are

# stored in the inode in the flag bits.


Use “ls -l” to see the permissions.

```
-rw-rw-r-- 1 grr 150 Aug 29 1995 calendar
-rw-r--r-- 1 grr 975 May 25 1999 device
-rwxrwxr-x 1 grr 5924 Jul 9 10:48 chars
-rw-rw-r-- 1 grr 124 Jul 9 10:47 chars.c
drwxr-sr-x 10 grr 512 Oct 14 1998 contools
drwxr-sr-x 9 grr 512 Oct 8 1998 contools-new
```

# Permission Bits

---

The permissions are grouped into three groups: User, Group, and Others.



# Permission Bits

---

- # To change the permissions of a file use the command `chmod`.

`chmod <u|g|o><+|-><r|w|x> filename`

Where

`<u|g|o>` is the owner, group or others.

`<+|->` Is to add or remove permissions

`<r|w|x>` Are read, write, execute permissions.

---

# Permission Bits Example

---

- # Make file “hello.txt” readable and writable by user and group but only readable by others

```
chmod u+rw hello.txt
```

```
chmod g+rw hello.txt
```

```
chmod o+r hello.txt
```

```
chmod o-w
```

```
hello.txt
```

□ Scripts and  
Executable files  
should have the

executable bit set

---

# Permission Bits

# Also you can change the permission bits all

at once using the **USER GROUP OTHERS** presentation in octal

<b>RWX</b>	<b>RWX</b>	<b>RWX</b>
------------	------------	------------

<b>110</b>	<b>110</b>	<b>100</b>	<b>- Binary</b>
------------	------------	------------	-----------------

<b>6</b>	<b>6</b>	<b>4</b>	<b>- Octal digits</b>
----------	----------	----------	-----------------------

**chmod 664 hello.c**

"r" corresponds with 4, "w" with 2, and "x" with 1.

# Process' Properties

---

# A process has the following properties:

- PID: Index in process table
  - Command and Arguments
  - Environment Variables
  - Current Dir
  - Owner (User ID)
  - Stdin/Stdout/Stderr
-

# Process ID

---

- # Uniquely identifies the processes among all live processes.
- # The initial process (init process) has ID of 0.
- # The OS assigns the numbers in ascending order.
- # The numbers wrap around when they reach the maximum and then are reused as long as there is no live process with the same processID.
- # You can programmatically get the process id with
  - `int getpid();`



# Command and Arguments

---

- # Every process also has a command that is executing (the program file or script) and 0 or more arguments.
  - # The arguments are passed to main.  
`int main(int argc, char **argv);`
  - # Argc contains the number of arguments including the command name.
  - # Argv[0] contains the name of the command
-

# Printing the Arguments

---

```
printargs.c:
int main(int argc, char **argv) {
    int i;
    for (i=0; i<argc; i++) {
        printf("argv[%d]=\"%s\"\n", i,
            argv[i]);
    }
}
```

```
gcc -o printargs printargs.c
./printargs hello world
argv[0]="./printargs"
argv[1]="hello"
argv[2]="world"
```

---

# Environment Variables

---

- # It is an array of strings that is inherited from the parent process.

Some important variables are:

- # ■ `PATH=/usr/bin` Stores the list of directories that contain commands to execute.
  - `USER=<login>` Contains the name of the user
  - `HOME=/homes/mzhu` Contains the home directory.
  - You can add Environment variables settings in `.bashrc` which is a Bash shell script that Bash runs whenever it is started interactively. It initializes an interactive shell session.
-

# Environment Variables

---

- # To set a variable from a shell use

`export A=B`

- Modify the environment globally. All processes called will get this change

`A=B`

- Modify environment locally. Only current shell process will get this change.

- # Example: Add a new directory to PATH

`export PATH=$PATH:/newdir`

---

# Printing Environment

# To print environment from a shell type  
`env`



```
mz10 — -bash — 80x24
Last login: Fri Feb  5 09:54:05 on console
CORE0782:~ mz10$ env
TERM_PROGRAM=Apple_Terminal
SHELL=/bin/bash
TERM=xterm-256color
TMPDIR=/var/folders/2p/_y68k3jj28g_7rdkgv0srf58_sbrs2/T/
Apple_PubSub_Socket_Render=/private/tmp/com.apple.launchd.sSEm5l6NR4/Render
TERM_PROGRAM_VERSION=404.1
TERM_SESSION_ID=8E43B928-04F4-49A0-B868-070DEB0064F2
USER=mz10
SSH_AUTH_SOCK=/private/tmp/com.apple.launchd.8h769qAd7f/Listeners
PATH=/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin
PWD=/Users/mz10
LANG=en_US.UTF-8
XPC_FLAGS=0x0
XPC_SERVICE_NAME=0
SHLVL=1
HOME=/Users/mz10
LOGNAME=mz10
SECURITYSESSIONID=186a7
_=/usr/bin/env
CORE0782:~ mz10$
```

# Current Directory

---

- # Every process also has a current directory.
- # The open file operations such as `open()` and `fopen()` will use the current directory to resolve relative paths.
- # If the path does not start with “/” then a path is relative to the current directory.

`/etc/hello.c` – Absolute

`path hello.c` – Relative  
path.

To change the directory use “`cd dir`” in a shell  
or

`chdir(dir)` inside a program

---

# Process' User ID

---

- # A process always runs in behalf of a user represented by the User ID.

- # The UID is inherited from the parent process.

- # Only root can change the UID of a process at runtime using the “setuid(uid);” call.

- # This happens at login time. The login program runs as root but then after identifying the user, the OS switches the UID to that user and runs the shell.

- # It also can be done with the comand “su user” that prompts for that user’s password.

- # Also “sudo user command” runs the command as



# Stdin/Stdout/Stderr

---

- # Also a process inherits from the parent a stdin/stdout and stderr.
- # They are usually the keyboard and the terminal but they can be redirected.



# PIPES

---

# In UNIX you can connect the output of a command to the input of another using PIPES.

The Pipe is a command in Linux that lets you use two or more commands such that output of one command serves as input to the next. In short, the output of each process directly as input to the next one like a pipeline. The symbol '|' denotes a pipe.

Example:

```
cat filename | less
```

pipe the output of the 'cat' command to 'less' which will show you only one scroll length of content at a time.

---



---

# Common UNIX Commands

---

# Common UNIX Commands

---

- # There are many UNIX commands that can be very useful
  - # The output of one can be connected to the input of the other using PIPES
  - # They are part of the UNIX distribution.
  - # Open source implementations exist and they are available in Linux.
-

# ls – list files and directories

# `ls <options> file list`

To show a long format list of files (permissions, ownership, size, and modification date), use:

# Examples

```
ls -l [path]
```

List all files, including important dotfiles (hidden files), use:

`ls -al`

```
ls -a [path]
```

- Lists all files including hidden files (files that start with “.”) and timestamps.

`ls -R dir`

- Lists recursively all directories and subdirectories in dir.

# mkdir – Make a directory

---

# mkdir <options> dir1 ...

# Examples

mkdir dir1

Create directory dir1

mkdir -p dir1/dir2/dir3

Also make parent directory if they do not exist.

---

# cp – Copy files

---

# `cp <options> file1 file2 ... destdir`

Copies one or more files to a destination.

# Examples:

`cp a.txt dir1`

Copies file a.txt into dir1

`cp a.txt b.txt`

Create a copy of a.txt called b.txt

`cp -R dir1 dir2`

Copy recursively directories and subdirectories of dir1 into dir2.

---

# mv – Move a file

---

# mv file1 destdir

Moves file1 to destdir

# Examples:

mv a.txt dir1

Move file a.txt into directory dir1

mv a.txt b.txt

Rename file a.txt into b.txt

---

# rm - Remove a file

---

# `rm <options> file1 file2 ...`

Removes a list of files

# Examples:

`rm a.txt b.txt`

Remove files a.txt and b.txt

`rm -R dir1`

Remove dir1 and all its contents.

---



# grep – Find lines

---

# `grep <options> pattern file1 file2 ...`

■ Print lines that contain “pattern”

# Examples:

`grep hello a.txt`

Print the lines in a.txt that contain hello

---

# man – Print manual pages

---

# man <options> command

- Print the manual page related to command.

# Examples:

man cp

Print the manual pages related to copy

man -k pthread

Print all manual pages that contain  
string “pthread”.

man -s 3 exec

# Print manual page of exec from  
section 3

Man Pages are divided into  
sections:

Section 1 – UNIX commands ( E.g. cp, mv  
etc) Section 2 – System Calls (E.g. fork)

Section 3 – C Standard Library

Example “man -s 1 printf” and “man -s 3c  
printf” give different man page

# whereis - Where a file is located

---

## # Where file

- Prints the path of where a file is located.
- It only works if the OS created a database with the files in the file system.
- Example  
bash-

# which – Path of a command

# which command

- Prints the path of the **command** that will be executed by the shell if

```
“CORE0782:~ mz10$ which ps
/bin/ps
CORE0782:~ mz10$ ps
ps: list PID TTY          TIME CMD
processe 18647 ttys000    0:00.09 -bash
CPU time  19603 ttys000    0:00.02 man sudo
          19604 ttys000    0:00.01 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19605 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19608 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19609 ttys000    0:00.01 /usr/bin/less -is
          19614 ttys000    0:00.01 man su
          19615 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19616 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19619 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          19620 ttys000    0:00.01 /usr/bin/less -is
          72312 ttys000    0:00.02 man cp
          72313 ttys000    0:00.01 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          72314 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          72317 ttys000    0:00.00 sh -c (cd '/usr/share/man' && /usr/bin/tbl '/usr/share
          72318 ttys000    0:00.02 /usr/bin/less -is
```

# head – Lists the first few lines

---

## # head file

List the first 10 lines of a file

### Example:

```
head myprog.c
```

Lists the first 10 lines of myprog.c

```
head -5 myprog.c
```

List the first 5 lines of myprog.

---

# tail – Lists the last few lines

---

## # tail file

- List the last 10 lines of a file

## # Example:

`tail myprog.c`

list the last 10 lines of myprog.c

`tail -3 myprog.c`

list the last 3 lines of myprog.c

`tail -f a.log`

It will periodically print the lines of a.log as they are added.

---

# find - Execute a command for a group of files

---

# find – Recursively descends the directory hierarchy for each path seeking files that match a Boolean expression.

# Examples:

```
find . -name "*.conf"
```

Find recursively all files with .conf extension.

```
find /home/mz10 -name "*.conf"
```

Find the file in the specific directory

---